

Hands-On with Heat – OpenStack Service Orchestration

Useful Information

The address for your cloud is <http://192.168.124.81>

Navigate to Project → Orchestration → Stacks and click the button to create a new stack

You may use one of the following methods for inputting your template:

1. Select “Direct Input” in the stack creation pop-up and either type or copy & paste the template from this lab guide into the appropriate box.
2. Copy the contents of the template to a text file, save it to your desktop, and then select “File” as your template source and upload the file.

Delete your stack after each lab exercise.

REMEMBER – INDENTATION/SPACING/CAPITALIZATION MATTERS IN YAML

Exercise 1 – “Hello World”

1. Using one of the steps above, input the text between the dotted lines into the stack creation dialogue for the 'Template Data' field. Leave every other field as is then hit 'Next'.

```
-----  
heat_template_version: 2013-05-23  
  
description: Simple template to deploy a single compute instance  
  
resources:  
  blog:  
    type: OS::Nova::Server  
    properties:  
      key_name: rashford  
      image: Wordpress-0.0.8-kvm  
      flavor: m1.small  
-----
```

Exercise 2 – “Parameters”

Using one of the steps above, input the text between the dotted lines into the stack creation dialogue

```
-----
heat_template_version: 2013-05-23

description: Simple template to deploy a single compute instance

parameters:
  key_name:
    type: string
    label: Key Name
    description: Name of key-pair to be used for compute instance
  image_id:
    type: string
    label: Image ID
    description: Image to be used for compute instance
  instance_type:
    type: string
    label: Instance Type
    description: Type of instance (flavor) to be used

resources:
  my_instance:
    type: OS::Nova::Server
    properties:
      key_name: { get_param: key_name }
      image: { get_param: image_id }
      flavor: { get_param: instance_type }
-----
```

Valid inputs to launch your parameters template:

- Key Name = rashford
- Image ID = Wordpress-0.0.8-kvm
- Instance Type = m1.small

Exercise 3 – “Defaults”

Using the template text below as a guide, create a **NEW** stack and **MODIFY** the “Parameters” lab template above to include default input for **both** parameters

```
-----
instance_type:
  type: string
  label: Instance Type
  description: Type of instance (flavor) to be used
  default: m1.small
-----
```

Exercise 4 – “Restricting User Input”

Using the template text below as a guide, **MODIFY** the “Parameters” lab template above to restrict the input for **both** parameters

```
-----
parameters:
  instance_type:
    type: string
    label: Instance Type
    description: Type of instance (flavor) to be used
    default: m1.small
    constraints:
      - allowed_values: [ m1.tiny, m1.small, m1.medium]
        description: Value must be one of m1.tiny, m1.small, m1.medium
-----
```

Exercise 5 – “2 Servers”

Using the template text below as a guide, **MODIFY** the “Parameters” lab template above to add in a second server

```
-----
database:
  type: OS::Nova::Server
  properties:
    key_name: rashford
    image: MySQL-0.0.3-kvm
    flavor: m1.small
-----
```

Exercise 6 – “Network”

Using the template text below, launch the blog service

```
-----
heat_template_version: 2013-05-23

description: Simple template to deploy a single compute instance

parameters:
  private_net_cidr:
    type: string
    description: Data network address (CIDR notation)
    default: 172.16.0.0/24

resources:
  wordpress:
    type: OS::Nova::Server
    properties:
      key_name: rashford
      image: Wordpress-0.0.8-kvm
      flavor: m1.small
    networks:
      - port: { get_resource: wordpress_data_port }
      - port: { get_resource: wordpress_port }
```

```
database:
  type: OS::Nova::Server
  properties:
    key_name: rashford
    image: MySQL-0.0.3-kvm
    flavor: m1.small
    networks:
      - port: { get_resource: mysql_data_port }

private_net:
  type: OS::Neutron::Net
  properties:
    name: data

private_subnet:
  type: OS::Neutron::Subnet
  properties:
    network_id: { get_resource: private_net }
    cidr: { get_param: private_net_cidr }

wordpress_port:
  type: OS::Neutron::Port
  properties:
    network_id: 79a090d5-b626-4bdd-9c17-84a6d01ff0b0
    security_groups: [ { get_resource: www_group } ]

wordpress_data_port:
  type: OS::Neutron::Port
  properties:
    network_id: { get_resource: private_net }
    fixed_ips:
      - subnet_id: { get_resource: private_subnet }

wordpress_floating_ip:
  type: OS::Neutron::FloatingIP
  properties:
    floating_network_id: 9c7c13ac-0e8e-40f3-b614-84d1de6ee672
    port_id: { get_resource: wordpress_port }

mysql_data_port:
  type: OS::Neutron::Port
  properties:
    network_id: { get_resource: private_net }
    security_groups: [ { get_resource: mysql_group } ]
    fixed_ips:
      - subnet_id: { get_resource: private_subnet }

www_group:
  type: OS::Neutron::SecurityGroup
  properties:
    description: Add security group rules for server
    name: www
    rules:
```

- remote_ip_prefix: 0.0.0.0/0
protocol: tcp
port_range_min: 80
port_range_max: 80
- remote_ip_prefix: 0.0.0.0/0
protocol: tcp
port_range_min: 443
port_range_max: 443
- remote_ip_prefix: 0.0.0.0/0
protocol: icmp

mysql_group:
type: OS::Neutron::SecurityGroup
properties:
description: Add security group rules for server
name: mysql
rules:
- remote_ip_prefix: 0.0.0.0/0
protocol: tcp
port_range_min: 3306
port_range_max: 3306
- remote_ip_prefix: 0.0.0.0/0
protocol: icmp

Exercise 7 – “Volumes”

MODIFY the “Network” lab template above to add in volume storage

wordpress_vol:
type: OS::Cinder::Volume
properties:
size: { get_param: wordpress_vol_size }

wordpress_vol_att:
type: OS::Cinder::VolumeAttachment
properties:
instance_uuid: { get_resource: wordpress }
volume_id: { get_resource: wordpress_vol }
mountpoint: /dev/vdb

mysql_vol:
type: OS::Cinder::Volume
properties:
size: { get_param: mysql_vol_size }

mysql_vol_att:
type: OS::Cinder::VolumeAttachment
properties:
instance_uuid: { get_resource: database }
volume_id: { get_resource: mysql_vol }
mountpoint: /dev/vdb

Exercise 8 - "Auto-Scaling/Load Balancing"

Pull the apache-autoscaling.yaml file and the lb_server.yaml file from the web server at <http://192.168.124.1/heat> and study them.

Launch the apache-autoscaling.yaml service and see it come up.

Experiment with different settings.

For example, how can you change the default number of servers that start up?

The current setup will deploy autoscale workloads one at a time. Can you set it to do it 3 at a time?

It currently is set to scale up if the cpu usage goes over 50% for 1 minute. Can you change these to make it scale up more aggressively?
